# Highlights

## EEG Data Augmentation with ERD-Guided and Class-Discriminative Diffusion Model

Chuncheng Liao, Wuqing Yu, Xiangcun Wang, Jiacai Zhang

- We propose an ERD-guided and class-discriminative conditional diffusion model that integrates classifier guidance and physiological constraints for high-quality EEG data augmentation.
- We introduce ERD laterality loss and frequency-aware spectral loss to ensure generated samples preserve physiologically plausible motor imagery patterns.
- We demonstrate significant improvements in classification accuracy across multiple scenarios, with 5.90% improvement in cross-session scenarios on BCI2a.
- Comprehensive experiments validate the method's superiority over traditional augmentation methods and other generative models across two datasets.

# EEG Data Augmentation with ERD-Guided and Class-Discriminative Diffusion Model

Chuncheng Liao[a], Wuqing Yu[a], Xiangcun Wang[a] and Jiacai Zhang[a,*]

[a]School of Artificial Intelligence, Beijing Normal University, No.19, Xinjiekouwai Street, Beijing, 100875, , China

ARTICLE INFO

*Keywords*:
EEG
data augmentation
diffusion model
classifier guidance
motor imagery

ABSTRACT

Motor Imagery Brain-Computer Interface (MI-BCI) systems are constrained by insufficient training samples and inter-subject variability, which limits classification accuracy and cross-subject generalization capability. To address this issue, this study proposes an ERD-guided and class-discriminative conditional diffusion probability model for high-quality electroencephalogram (EEG) data augmentation.Building upon the standard conditional diffusion probability model (Cond-DDPM), the proposed architecture incorporates classifier guidance from a pre-trained EEGNet teacher network to extract class-discriminative features, and integrates ERD laterality and frequency-aware spectral constraints as physiologically grounded guidance during the generation process. The study employs BCI Competition IV Datasets 2a and 2b, benchmarking against traditional (Gaussian Noise, SMOTE) and generative (VAE, GAN) methods. On BCI2a, our ERD-guided DDPM boosts cross-session accuracy by +5.90% (63.39%→69.29%), within-subject by +5.17% (72.61%→77.78%), and cross-subject by +1.66% (50.77%→52.43%); on BCI2b, it raises within-subject accuracy by +3.28% (83.63%→86.91%), cross-session by +1.15% (84.11%→85.26%), and cross-subject by +0.97% (71.18%→72.15%). Outperforming all baselines, our method generates high-fidelity, class-discriminative EEG data, addressing MI-BCI data scarcity.

## 1. Introduction

Electroencephalogram (EEG), as a direct representation of neural activity in the brain, holds significant application value in brain-computer interfaces (BCIs), clinical diagnosis, and neuroscience Nicolas-Alonso and Gomez-Gil (2012); Yuan and He (2014); Miller, Schalk, Fetz, den Nijs, Ojemann and Rao (2010). However, practically recorded EEG signals are susceptible to interference such as power-line noise and myoelectric artefacts, leading to degraded data quality and blurred features that severely impact downstream task performance Wolpaw, Birbaumer, McFarland, Pfurtscheller and Vaughan (2002); Mouawad, Dousty and Sattar (2018). Traditional data augmentation methods (e.g., linear filtering, wavelet transform) rely on prior assumptions, struggle to adapt to the non-stationarity, high dimensionality, and inter-individual variability of EEG, and may damage critical spatiotemporal characteristics Pfurtscheller and Neuper (2001); Schalk, McFarland, Hinterberger, Birbaumer and Wolpaw (2004); Xu, Xu, Ke, An, Liu and Ming (2020); Albuquerque, Gama and Figueiredo (2014), necessitating novel data enhancement approaches. In parallel, recent work has explored spiking neural network based decoders for intra-cortical brain signals Yang, Fu, Zhang, Zhang, Li and Wu (2026), highlighting the broader potential of advanced neural architectures for neural decoding in BCI systems.

Motor imagery (MI)-based BCI systems have garnered considerable attention due to their non-invasive nature Wolpaw et al. (2002); Pfurtscheller and Neuper (2001). Their core mechanism lies in the event-related desynchronization (ERD) and synchronization (ERS) phenomena within the sensorimotor cortex's $\mu$ band (8–13 Hz) and $\beta$ band (13–30 Hz) during MI tasks Pfurtscheller and Neuper (1997); Neuper and Pfurtscheller (2001), serving as reliable biomarkers for MI-BCI systems. Despite technological advancements, three major challenges persist: scarce training data constrains robust model development; substantial inter-subject variability impedes cross-subject generalization; and conventional augmentation methods fail to preserve key physiological features like ERD/ERS, potentially introducing detrimental artefacts. Recent work on cross-subject EEG emotion recognition has proposed spatiotemporal cross-brain region interaction networks to better capture subject-invariant representations An, Hu, Liu,

*Corresponding author

liaocc3450@bnu.edu.cn (C. Liao); 202431081038@mail.bnu.edu.cn (W. Yu); wangxiangcun@mail.bnu.edu.cn (X. Wang); jiacai.zhang@bnu.edu.cn (J. Zhang)

ORCID(s):

Gu, Zhang and Zhang (2025), and multi-view contrastive networks have been developed to enhance MI classification performance by exploiting complementary views of EEG representations Wang, Li, Chen and Wu (2025a), further underscoring the importance of modeling cross-subject variability and multi-view structure in neural decoding systems.

In recent years, deep learning–driven data augmentation methods have offered new insights for EEG signal processing, yet limitations have emerged during their evolution:

Unconditional generative models (e.g., DCGAN, $\beta$-VAE) can generate samples resembling real signal distributions but lack conditional information, resulting in low relevance to downstream tasks. Single-scale modeling overlooks EEG's multi-level spatiotemporal hierarchies Yoo, Uh, Chun, Kang and Ha (2019); Song, Meng and Ermon (2020).

Conditional generative models (e.g., Conditional GAN, Conditional $\beta$-VAE) improve task relevance through class labels but remain constrained by fixed-scale feature extraction. They fail to resolve dynamic correlations between spectral features and temporal resolution, particularly struggling to capture coupling patterns between low-frequency global trends and high-frequency local details Li, Wang, Zhang, Cui, Li, Zhou and Yang (2020); Wang, She and Ward (2021).

Conditional diffusion models (CDMs) (e.g., DDPM, EEG-CDM) preserve data distribution via reverse denoising and support conditioning. However, their traditional architecture depends on fixed-scale feature extractors (e.g., standard CNNs), and noise scheduling strategies neglect EEG's time-frequency properties, causing high-frequency detail loss or low-frequency trend distortion Liu, Zhou, Yuan and Chen (2018); Chen, Zhang, Zen, Weiss, Norouzi and Chan (2021). Although subsequent studies attempt multi-scale improvements (e.g., frequency band division, hierarchical diffusion models), predefined bands and insufficient cross-scale coupling modeling limit adaptation to EEG's dynamic multi-scale characteristics Huang, Liu, Van Der Maaten and Weinberger (2017); Zhang, Wang, Yang, Sanford, Harmon, Turkbey, Roth, Xu, Xu, Ziyue et al. (2020); Wang and Perez (2017). Recent work has begun to apply diffusion probabilistic models directly to EEG synthetic data generation, demonstrating the potential of this family of models for BCI-oriented data augmentation Tosato, Dalbagno and Fumagalli (2023); Torma and Szegletes (2025); Wang, Li, Chen and Wu (2025b); beyond augmentation, diffusion-based frameworks have also been leveraged to reconstruct visual images from EEG recordings Zeng, Xia, Qian, Hattori, Wang and Kong (2023), highlighting the broader applicability of diffusion models for EEG-driven generative tasks.

Critically, existing methods fail to generate samples with strong class-discriminative features, yielding statistically plausible data that lack the semantic clarity needed for effective classification. To address these issues, this paper proposes an ERD-guided and class-discriminative conditional diffusion model, aiming to achieve high-quality EEG data augmentation by incorporating classifier guidance from a pre-trained teacher network and ERD-based physiological constraints into the diffusion process. Core innovations include:

1. We integrate gradient-based guidance from a pre-trained EEGNet classifier to steer the generation process toward producing samples with clear class-specific features. This ensures that generated data maintains semantic correctness for specific motor intentions while preserving the underlying data distribution.
2. We enforce physiologically plausible contralateral hemispheric desynchronization patterns by penalizing deviations from class-specific alpha-band power asymmetry between C3 and C4 electrodes. This ERD-guided constraint encourages generated samples to exhibit realistic motor imagery-related ERD patterns.
3. We employ a spectral preservation loss that emphasizes physiologically relevant frequency bands (alpha: 8-13 Hz, delta: 1-4 Hz) through weighted power spectral density matching. This ERD-related spectral guidance encourages the model to preserve critical neural rhythms associated with motor imagery tasks.
4. Downstream classification tasks are embedded into the diffusion training process through classifier guidance. Gradient backpropagation aligns generated data directly with classification performance improvement, mitigating the disconnect between "generation" and "application" stages in conventional pipelines.

Experiments on the BCI Competition IV Dataset 2a demonstrate that our ERD-guided and class-discriminative DDPM generates high-quality samples that preserve spectral characteristics while enhancing class separability. Classification performance improves across within-subject, cross-session, and cross-subject scenarios, with the most significant improvement in cross-session (+5.90%). Additional validation on the BCI Competition IV Dataset 2b confirms the generalizability of our method, with consistent improvements observed across all three scenarios. Comprehensive comparisons against traditional methods (Gaussian Noise, SMOTE) and other generative models (GAN, VAE) validate the superiority of our approach for EEG data augmentation across diverse dataset configurations.

The remainder of this paper is organized as follows: Section 2 introduces the neurophysiological mechanisms of motor imagery and the fundamentals of the diffusion probability model, and elaborates on the ERD-guided and class-discriminative diffusion model methodology; Section 3 describes the experimental setup and datasets, and presents the experimental results; Section 4 discusses the research findings; and Section 5 concludes the paper and outlines future work.

## 2. Materials And Methods

### 2.1. Neurophysiological Mechanisms of Motor Imagery

#### *2.1.1. ERD/ERS Phenomena*

During motor imagery, characteristic rhythmic modulations occur in the sensorimotor cortex, primarily manifested as: Event-Related Desynchronization (ERD)—during motor imagination or execution, power reduction in the $\mu$ band (8–13 Hz) and $\beta$ band (13–30 Hz) over specific cortical regions reflects a transition from synchronized to desynchronized neuronal population activity, indicating cortical activation Pfurtscheller and Neuper (1997). Event-Related Synchronization (ERS)—post-movement or in non-relevant brain regions, increased rhythmic power signifies restored neuronal synchronization, reflecting cortical inhibition or disengagement Neuper and Pfurtscheller (2001).

#### *2.1.2. Frequency Band Specificity*

The $\mu$ rhythm (8–13 Hz) originates primarily from the sensorimotor cortex (near electrodes C3, Cz, and C4). During motor imagery, it exhibits contralateral hemispheric ERD: for example, when imagining right-hand movement, $\mu$-band power decreases in the left hemisphere (C3). The $\beta$ rhythm (13–30 Hz), also generated in the sensorimotor cortex, displays ERD/ERS patterns similar to but more complex than the $\mu$ rhythm. Notably, $\beta$-band ERD emerges during movement preparation, followed by a post-movement beta rebound Pfurtscheller and Neuper (2001); McFarland, Miner, Vaughan and Wolpaw (2000).

#### *2.1.3. Spatial Specificity*

Motor imagery of different body parts activates distinct cortical regions: hand movements engage the contralateral central regions (C3/C4); foot movements activate the central-parietal area (Cz and surrounding electrodes); and tongue movements involve the bilateral central regions Schirrmeister, Springenberg, Fiederer, Glasstetter, Eggensperger, Tangermann, Hutter, Burgard and Ball (2017); Ang, Chin, Zhang and Guan (2008). This spatially, spectrally, and temporally specific pattern serves as a critical basis for classification in MI-BCI systems.

### 2.2. ERD-Guided and Class-Discriminative Diffusion Model Overview

Figure 1 illustrates the architecture of the proposed ERD-guided and class-discriminative diffusion model. The framework consists of two main components: (1) a conditional DDPM backbone for generative modeling, and (2) a pre-trained EEGNet classifier providing class-discriminative guidance during the generation process, jointly regularized by ERD laterality and spectral preservation losses.

### 2.3. Conditional Diffusion Probabilistic Model

The conditional DDPM backbone employs a multi-scale conditional U-Net-based noise prediction network Ho, Jain and Abbeel (2020); Nichol and Dhariwal (2021), taking as input the noisy EEG signals $\mathbf{x}_t$, the time step $t$, and the class label $c$. The network architecture consists of three downsampling stages and three upsampling stages, with residual connections and self-attention mechanisms incorporated. The architecture employs an enhanced class embedding module that uses a multi-layer MLP with residual connections to encode class labels into a 256-dimensional embedding space. Time step information is encoded using sinusoidal positional embeddings and projected through a two-layer MLP to match the class embedding dimension. The combined time and class embeddings are integrated at each encoder and decoder stage via condition-aware residual blocks that apply linear projections of the condition vector as additive biases, similar to Feature-wise Linear Modulation (FiLM) layers.

#### *2.3.1. Forward Diffusion Process*

The forward process gradually adds Gaussian noise to real EEG data $\mathbf{x}_0$ over $T$ timesteps, producing a sequence of increasingly noisy latent variables $\mathbf{x}_1, \ldots, \mathbf{x}_T$:

$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1-\beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I}) \tag{1}$$

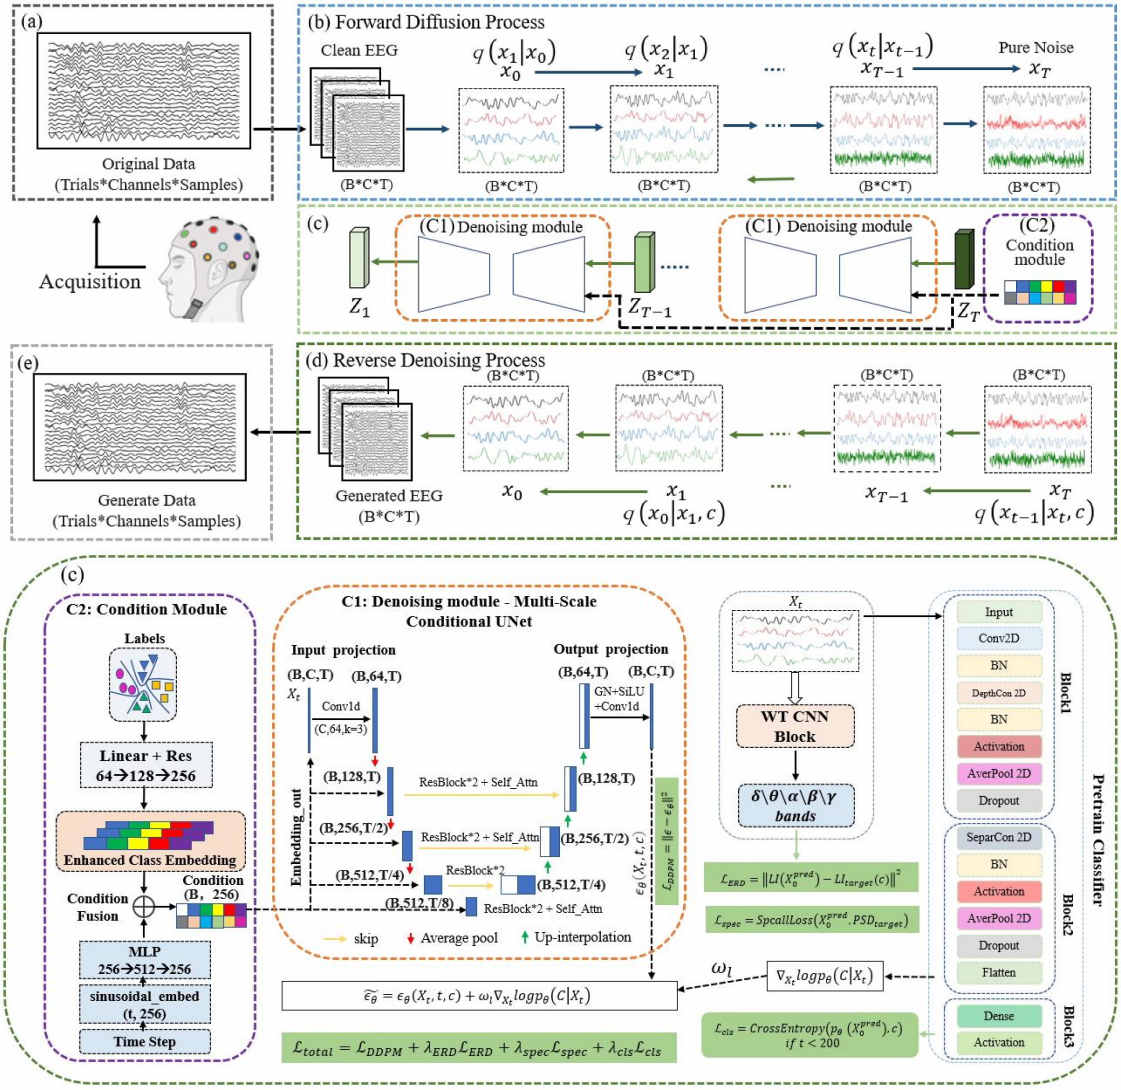


**Figure 1:** ERD-guided and class-discriminative diffusion model architecture.(a) Raw EEG data, with a shape of (trials, channels, samples).(b) Forward diffusion process. Starting from the raw EEG data $\mathbf{x}_0$, Gaussian noise is progressively added over T=1000 time steps according to a preset noise schedule. This process transforms the original clear EEG signal into a state of pure random noise, laying the foundation for the model to learn the reverse process.(c) Model backbone. Integrating a conditional DDPM, it combines classifier-based guidance with ERD and spectral constraints. This aims to ensure that the generated samples possess distinct class-specific features while retaining physiologically plausible signal properties.(d) Reverse denoising process. Initiating from the pure noise state, this process utilizes the pretrained backbone network and incorporates conditional guidance mechanisms. By progressively predicting and removing noise, it accurately reconstructs EEG data that exhibits both physiological authenticity and class-discriminative features.(e) Generated EEG data,with a shape of (trials, channels, samples).

where $\beta_t$ is the noise schedule controlling the variance at step $t$. Using the reparameterization trick, we can directly sample $\mathbf{x}_t$ from $\mathbf{x}_0$:

$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1-\bar{\alpha}_t}\epsilon, \quad \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \tag{2}$$

where $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^{t} \alpha_s$.

#### 2.3.2. Reverse Denoising Process

The reverse process learns to denoise $\mathbf{x}_t$ back to $\mathbf{x}_0$, conditioned on class label $c$:

$$p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, c) = \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t, c), \boldsymbol{\Sigma}_\theta(\mathbf{x}_t, t, c)) \tag{3}$$

A multi-scale conditional U-Net neural network $\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t, c)$ is trained to predict the noise $\boldsymbol{\epsilon}$ added at timestep $t$, with the training objective:

$$\mathcal{L}_{\text{DDPM}} = \mathbb{E}_{\mathbf{x}_0,t,\boldsymbol{\epsilon}}[\|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t, c)\|^2] \tag{4}$$

#### 2.3.3. Multi-Scale Conditional U-Net Architecture

The denoising network follows a U-Net architecture with the following structure:

- **Input Projection**: A 3×1 convolutional layer projects the input from $C$ channels (e.g., 22 for BCI2a) to 64 feature channels.
- **Encoder Path**: Three downsampling stages progressively reduce temporal resolution while increasing feature channels:
  - **Stage 1**: Two residual blocks (64 → 128 channels) followed by a self-attention block, maintaining full temporal resolution.
  - **Stage 2**: Average pooling (2× downsampling), two residual blocks (128 → 256 channels), and a self-attention block.
  - **Stage 3**: Average pooling (2× downsampling), two residual blocks (256 → 512 channels).
- **Bottleneck**: Two residual blocks (512 channels) with a self-attention block in between, operating at the lowest temporal resolution (T/8).
- **Decoder Path**: Three upsampling stages progressively restore temporal resolution while decreasing feature channels:
  - **Stage 3**: Linear interpolation upsampling, concatenation with encoder skip connections, two residual blocks (1024 → 512 → 256 channels).
  - **Stage 2**: Linear interpolation upsampling, concatenation with skip connections, two residual blocks (512 → 256 → 128 channels), and a self-attention block.
  - **Stage 1**: Linear interpolation upsampling, concatenation with skip connections, two residual blocks (256 → 128 → 64 channels), and a self-attention block.
- **Output Projection**: Group normalization, SiLU activation, and a 3×1 convolutional layer projects features back to $C$ channels.

Each residual block consists of two 3×1 convolutional layers with GroupNorm (8 groups) and SiLU activations, with condition injection via linear projection of the combined time and class embedding as an additive bias. Self-attention blocks enable the network to capture long-range temporal dependencies in EEG signals, which is crucial for modeling ERD/ERS patterns that span multiple time points.

#### 2.3.4. Enhanced Class Embedding

The class embedding module employs a multi-layer MLP architecture with residual connections:

- **Layer 1**: Embedding lookup maps class indices to 64-dimensional vectors.
- **Layer 2**: Linear projection (64 → 128) with SiLU activation.
- **Layer 3**: Linear projection (128 → 256) with SiLU activation.
- **Layer 4**: Linear projection (256 → 256) with residual connection from Layer 2 (projected to 256 dimensions).

This enhanced embedding architecture provides richer class representations compared to simple embedding lookups, improving the model's ability to distinguish between different motor imagery classes.

## 2.4. Classifier Guidance Mechanism

### 2.4.1. Classifier Guidance

To ensure generated samples maintain semantic correctness and strong class-discriminative features, we employ classifier guidance Ho and Salimans (2022). A pre-trained EEGNet classifier $p_\phi(c|\mathbf{x}_t)$ provides gradients that steer generation toward the target class:

$$\nabla_{\mathbf{x}_t} \log p_\phi(c|\mathbf{x}_t) = \nabla_{\mathbf{x}_t} \text{logit}_c(\mathbf{x}_t) \tag{5}$$

where $\text{logit}_c(\mathbf{x}_t)$ denotes the classifier's logit output for class $c$.

### 2.4.2. ERD Laterality Loss

To encourage physiologically plausible motor imagery patterns, we incorporate an ERD laterality loss that enforces contralateral hemispheric desynchronization. The laterality index quantifies the power asymmetry between C3 and C4 electrodes in the alpha band (8-13 Hz):

$$\text{LI}(x) = \frac{P^\alpha_{\text{C4}} - P^\alpha_{\text{C3}}}{P^\alpha_{\text{C4}} + P^\alpha_{\text{C3}} + \epsilon} \tag{6}$$

where $P^\alpha_{\text{C3}}$ and $P^\alpha_{\text{C4}}$ denote alpha-band power at electrodes C3 and C4, respectively. For left-hand motor imagery, the laterality index should be positive (stronger C4 activity), while for right-hand imagery, it should be negative (stronger C3 activity) Pfurtscheller and Neuper (1997). The ERD loss penalizes deviations from class-specific target laterality values:

$$\mathcal{L}_{\text{ERD}} = \|\text{LI}(x_{\text{pred}}) - \text{LI}_{\text{target}}(c)\|^2 \tag{7}$$

### 2.4.3. Spectral Preservation Loss

To preserve physiologically relevant frequency characteristics, we incorporate a spectral loss that emphasizes critical EEG frequency bands. The spectral loss computes the mean squared error between the power spectral densities (PSD) of generated and target signals in log-space, with additional weighting for the alpha band (8-13 Hz) and delta band (1-4 Hz):

$$\mathcal{L}_{\text{spec}} = \|P_{\text{pred}} - P_{\text{target}}\|^2 + w_\alpha \|P^\alpha_{\text{pred}} - P^\alpha_{\text{target}}\|^2 + w_\delta \|P^\delta_{\text{pred}} - P^\delta_{\text{target}}\|^2 \tag{8}$$

where $P$ denotes log-transformed PSD, superscripts $\alpha$ and $\delta$ indicate frequency band masking, and $w_\alpha = 2.0$, $w_\delta = 1.5$ are band-specific weights that emphasize motor imagery-relevant rhythms Pfurtscheller and Neuper (1997).

### 2.4.4. Guided Generation

The final guided noise prediction combines the standard DDPM prediction with the classifier guidance signal:

$$\tilde{\epsilon}_\theta = \epsilon_\theta(\mathbf{x}_t, t, c) + w_l \nabla_{\mathbf{x}_t} \log p_\phi(c|\mathbf{x}_t) \tag{9}$$

where $\epsilon_\theta(\mathbf{x}_t, t, c)$ is the noise prediction from the multi-scale conditional U-Net conditioned on noisy input $\mathbf{x}_t$, timestep $t$, and class label $c$; $w_l$ is a weighting hyperparameter balancing the guidance strength. During training, the model optimizes a combined objective:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{DDPM}} + \lambda_{\text{ERD}}\mathcal{L}_{\text{ERD}} + \lambda_{\text{spec}}\mathcal{L}_{\text{spec}} + \lambda_{\text{cls}}\mathcal{L}_{\text{cls}} \tag{10}$$

where $\mathcal{L}_{\text{ERD}}$ is the ERD laterality loss, $\mathcal{L}_{\text{spec}}$ is the spectral preservation loss, $\mathcal{L}_{\text{cls}}$ is the classification loss applied at low noise levels, and $\lambda_{\text{ERD}}$, $\lambda_{\text{spec}}$, $\lambda_{\text{cls}}$ are loss weights.

Algorithm 1 summarizes the complete training procedure. During training, at each iteration, we sample a batch of real EEG data along with their class labels, then sample a random timestep $t$ and add Gaussian noise to obtain $\mathbf{x}_t$. The multi-scale conditional U-Net, incorporating enhanced class embeddings and self-attention mechanisms, predicts the added noise $\epsilon_\theta(\mathbf{x}_t, t, c)$, from which we can recover the predicted clean sample $\mathbf{x}_0^{\text{pred}}$. We then compute all four loss components: the noise prediction loss $\mathcal{L}_{\text{DDPM}}$ ensures accurate denoising; the ERD laterality loss $\mathcal{L}_{\text{ERD}}$ enforces physiologically plausible hemispheric asymmetry patterns; the spectral loss $\mathcal{L}_{\text{spec}}$ preserves frequency-domain

characteristics with emphasis on alpha and delta bands; and the classification loss $\mathcal{L}_{\text{cls}}$ (applied only when $t < 200$) enhances class discriminability. The total loss is a weighted combination of these four components, and the model parameters are updated via gradient descent.

Algorithm 2 describes the generation process during inference. Starting from pure Gaussian noise $\mathbf{x}_T$, we iteratively denoise through $T$ steps. At each step, the trained multi-scale conditional U-Net predicts the noise component $\epsilon_\theta(\mathbf{x}_t, t, c)$ conditioned on the target class $c$, and we apply classifier guidance by adding the gradient of the log-probability from the pre-trained classifier. This guidance steers the generation toward the target class while maintaining the learned data distribution. The guidance weight $w_l$ controls the strength of this steering effect—higher values produce samples with stronger class-specific features but may sacrifice diversity.

**Algorithm 1** ERD-guided And Class-Discriminative DDPM Training

**Input:** Training data $\mathcal{D} = \{(\mathbf{x}_0^{(i)}, c^{(i)})\}_{i=1}^N$, pre-trained classifier $p_\phi$, diffusion steps $T$, loss weights $\lambda_{\text{ERD}}$, $\lambda_{\text{spec}}$, $\lambda_{\text{cls}}$
**Output:** Trained denoising network $\theta$
**repeat**
  Sample batch $\{(\mathbf{x}_0, c)\}$ from training data
  Sample timestep $t \sim \text{Uniform}(1, T)$
  Sample noise $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
  Compute noisy sample: $\mathbf{x}_t = \sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1-\bar{\alpha}_t}\epsilon$
  Predict noise: $\epsilon_\theta \leftarrow$ Multi-Scale Cond. U-Net$(\mathbf{x}_t, t, c)$
  Predict clean sample: $\mathbf{x}_0^{\text{pred}} = (\mathbf{x}_t - \sqrt{1-\bar{\alpha}_t}\epsilon_\theta)/\sqrt{\bar{\alpha}_t}$
  Compute losses:
    $\mathcal{L}_{\text{DDPM}} = \|\epsilon - \epsilon_\theta\|^2$
    $\mathcal{L}_{\text{ERD}} = \|\text{LI}(\mathbf{x}_0^{\text{pred}}) - \text{LI}_{\text{target}}(c)\|^2$
    $\mathcal{L}_{\text{spec}} = \text{SpectralLoss}(\mathbf{x}_0^{\text{pred}}, \text{PSD}_{\text{target}})$
    $\mathcal{L}_{\text{cls}} = \text{CrossEntropy}(p_\phi(\mathbf{x}_0^{\text{pred}}), c)$ if $t < 200$
  Compute total loss:
    $\mathcal{L} = \mathcal{L}_{\text{DDPM}} + \lambda_{\text{ERD}}\mathcal{L}_{\text{ERD}} + \lambda_{\text{spec}}\mathcal{L}_{\text{spec}} + \lambda_{\text{cls}}\mathcal{L}_{\text{cls}}$
  Update $\theta$ using gradient descent: $\theta \leftarrow \theta - \eta\nabla_\theta\mathcal{L}$
**until** convergence

**Algorithm 2** ERD-guided And Class-Discriminative DDPM Generation (Inference)

**Input:** Target class $c$, guidance weight $w_l$, trained model $\epsilon_\theta$
**Output:** Generated EEG sample $\mathbf{x}_0$
Initialize $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
**for** $t = T, \ldots, 1$ **do**
  Predict noise: $\epsilon_\theta \leftarrow$ Multi-Scale Cond. U-Net$(\mathbf{x}_t, t, c)$
  Compute classifier gradient: $\nabla_{\mathbf{x}_t} \log p_\phi(c|\mathbf{x}_t)$
  Apply classifier guidance:
    $\tilde{\epsilon}_\theta \leftarrow \epsilon_\theta + w_l\nabla_{\mathbf{x}_t} \log p_\phi(c|\mathbf{x}_t)$
  Denoise: $\mathbf{x}_{t-1} \sim p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, \tilde{\epsilon}_\theta)$
**end for**
**return** $\mathbf{x}_0$

## 2.5. Evaluation Metrics

We evaluate the effectiveness of our data augmentation approach using both data quality metrics and classification performance metrics:

### *2.5.1. Data Quality Metrics*

To quantify the fidelity of generated EEG data, we compute frequency-domain similarity metrics on Subject 1 data:

**PSD Correlation**: Pearson correlation between the power spectral density (PSD) of generated and real data, measuring overall frequency-domain similarity Welch (1967).

**Frequency Similarity**: Mean correlation across five physiologically relevant frequency bands (delta: 1-4 Hz, theta: 4-8 Hz, alpha: 8-13 Hz, beta: 13-30 Hz, gamma: 30-50 Hz), quantifying band-specific preservation Pfurtscheller and Neuper (1997).

These metrics evaluate spectral fidelity of generated data. However, high quality metrics alone do not guarantee utility for downstream tasks—generated data must also preserve class-discriminative features to improve classification performance.

#### *2.5.2. Classification Performance Metrics*

The primary evaluation metric is classification accuracy achieved on the test set after augmenting the training set with generated data Lawhern, Solon, Waytowich, Gordon, Hung and Lance (2018); Craik, He and Contreras-Vidal (2019). This metric quantifies the performance improvement of downstream tasks facilitated by the synthetic data, reflecting both the information content of the generated samples and their semantic consistency with real data.

$$\text{Acc} = \frac{1}{N_{\text{test}}} \sum_{i=1}^{N_{\text{test}}} \mathbb{I}(\hat{y}_i = y_i) \tag{11}$$

where $N_{\text{test}}$ is the number of test samples, $\hat{y}_i$ is the predicted label, $y_i$ is the true label, and $\mathbb{I}(\cdot)$ is the indicator function. We report mean accuracy and standard deviation across all subjects for each experimental scenario.

#### *2.5.3. Statistical Significance Testing*

To assess whether the improvements achieved by our method are statistically significant, we conduct paired t-tests comparing DDPM-augmented results against the no-augmentation baseline across all subjects. We also compute Cohen's d effect size to quantify the magnitude of improvement:

$$d = \frac{\bar{x}_{\text{DDPM}} - \bar{x}_{\text{baseline}}}{s_{\text{pooled}}} \tag{12}$$

where $\bar{x}$ denotes mean accuracy and $s_{\text{pooled}}$ is the pooled standard deviation. Effect sizes are interpreted as small ($d \approx 0.2$), medium ($d \approx 0.5$), or large ($d \approx 0.8$) Cohen (1988).

## 3. Experiment And Results

### 3.1. Experimental Dataset

#### *3.1.1. BCI Competition IV Dataset 2a*

The BCI Competition IV Dataset 2a Brunner, Leeb, Müller-Putz, Schlögl and Pfurtscheller (2008) comprises EEG recordings from 9 subjects performing four motor imagery tasks: left hand, right hand, both feet, and tongue movement. Each subject completed 288 trials (72 per class) across two sessions recorded on different days. The data were acquired using 22 Ag/AgCl electrodes positioned according to the international 10-20 system, with a sampling rate of 250 Hz.

#### *3.1.2. BCI Competition IV Dataset 2b*

Similarly derived from 9 subjects, this dataset simplifies the task to binary classification of left- and right-hand motor imagery Tangermann, Müller, Aertsen, Birbaumer, Braun, Brunner, Leeb, Mehring, Miller, Müller-Putz, Nolte, Pfurtscheller, Preissl, Schalk, Schlögl, Vidaurre, Waldert and Blankertz (2012). Each subject performed 288 trials, with a signal acquisition time window of [3–7] seconds and a sampling rate range of 4–40 Hz. The data division assigns the first three sessions as the training set and the last two sessions as the test set, facilitating research on binary motor imagery tasks.

The experimental paradigm for each trial proceeded as follows: (1) a fixation cross appeared for 2 seconds, instructing subjects to prepare; (2) a visual cue (arrow) indicated the MI task for 1.25 seconds; (3) subjects performed continuous MI until the cross disappeared at t=6 seconds; (4) a short inter-trial interval allowed rest. For our experiments, we extracted EEG segments from 2-6 seconds post-cue, encompassing the active MI period when ERD phenomena are most pronounced.

## 3.2. Experimental Setup

### *3.2.1. Data Preprocessing*

EEG signals were preprocessed using a standardized pipeline implemented in Python with the MNE library Gramfort, Luessi, Larson, Engemann, Strohmeier, Brodbeck, Parkkonen and Hämäläinen (2013):

1. **Bandpass Filtering**: A 4th-order Butterworth bandpass filter (0.5-40 Hz) was applied to remove low-frequency drifts and high-frequency noise while preserving relevant MI-related activity Blankertz, Dornhege, Krauledat, Müller and Curio (2007).
2. **Common Average Reference (CAR)**: Spatial filtering via CAR reduced common-mode artifacts and enhanced signal-to-noise ratio McFarland et al. (2000).
3. **Artifact Rejection**: Trials with excessive amplitude (>100 $\mu$V) were automatically rejected to exclude movement artifacts and electrode issues Delorme and Makeig (2004).
4. **Exponential Moving Standardization**: Channel-wise standardization with exponential smoothing normalized signal amplitudes while preserving temporal dynamics Schirrmeister et al. (2017).

Following preprocessing, each 4-second trial yielded a data sample of shape (22, 1000), corresponding to 22 channels and 1000 time points. The processed dataset comprised 5,184 trials across 9 subjects.

### *3.2.2. Training Procedure*

Training proceeds in two stages: (1) Pre-train the EEGNet classifier Lawhern et al. (2018) on real data until convergence; (2) Train the conditional DDPM using the base objective $\mathcal{L}_{\text{DDPM}}$ for 300 epochs with batch size 128, learning rate 1e-4, and Adam optimizer Kingma and Welling (2013). The noise schedule uses a linear beta schedule from 0.0001 to 0.02 over 1000 diffusion steps Ho et al. (2020).

### *3.2.3. Generation Procedure*

During generation, we set guidance weight $w_l = 2.0$ based on grid search validation Cubuk, Zoph, Mane, Vasudevan and Le (2019). For each class and subject, we generate synthetic samples matching the original data distribution. The generation process uses 1000 diffusion steps with a linear beta schedule from 0.0001 to 0.02 Ho et al. (2020).

## 3.3. Baseline Methods

We benchmark our proposed ERD-guided and class-discriminative DDPM model against representative approaches across three categories: traditional augmentation methods, generative models, and conditional diffusion models. Below are the characteristics of each method in the context of EEG data augmentation:

**Gaussian Noise** Zhang, Cisse, Dauphin and Lopez-Paz (2017): A simple augmentation technique that adds zero-mean Gaussian noise with standard deviation $\sigma = 0.1$ to raw EEG signals. This method aims to improve model robustness by introducing controlled perturbations.

**SMOTE (Synthetic Minority Over-sampling Technique)** Chawla, Bowyer, Hall and Kegelmeyer (2002): A data augmentation method that generates synthetic samples through linear interpolation between existing samples and their k-nearest neighbors in the feature space. For EEG data, we apply SMOTE in the flattened signal space.

**VAE (Variational Autoencoder)** Kingma and Welling (2013); Song et al. (2020): A latent-variable generative model that learns a compressed latent representation of EEG data via an encoder network and synthesizes novel samples through a decoder network. We implement a conditional VAE with class labels to generate class-specific samples.

**GAN (Generative Adversarial Network)** Yoon, Jarrett and van der Schaar (2019); Hartmann, Schirrmeister and Ball (2018): A generative model based on adversarial training between a generator and discriminator. We implement a conditional GAN architecture specialized for time-series EEG data, incorporating temporal convolutions to preserve temporal dependencies.

**Our Method - ERD-Guided And Class-Discriminative DDPM**: Building upon conditional DDPM Ho et al. (2020); Nichol and Dhariwal (2021), our method employs a U-Net neural network to learn reverse denoising processes through a Markov chain structure. The key innovation is the integration of classifier guidance from a pre-trained EEGNet teacher network together with ERD laterality and spectral preservation losses, which jointly provide class-discriminative and physiologically grounded signals during the generation process to ensure that generated samples have clear class-specific features while preserving realistic ERD patterns.

**Table 1**
Data Quality Metrics Comparison on BCI2a Dataset (Subject 1 Session 1)

| Method | PSD Corr. | Freq. Sim. | Mean Corr. |
|---|---|---|---|
| Gaussian Noise | 0.787 | 0.519 | 0.653 |
| SMOTE | 0.787 | 0.529 | 0.658 |
| VAE | 0.779 | 0.429 | 0.604 |
| GAN | 0.742 | 0.271 | 0.507 |
| **DDPM (Ours)** | **0.948** | **0.637** | **0.793** |

**Table 2**
Band-wise PSD Correlations on BCI2a Dataset (Subject 1 Session 1)

| Method | Delta | Theta | Alpha | Beta | Gamma |
|---|---|---|---|---|---|
| Gaussian Noise | 0.964 | 0.096 | 0.279 | 0.351 | 0.904 |
| SMOTE | 0.964 | 0.110 | 0.285 | 0.382 | 0.904 |
| VAE | 0.722 | 0.077 | 0.186 | 0.348 | 0.815 |
| GAN | 0.247 | -0.132 | 0.133 | 0.398 | 0.709 |
| **DDPM (Ours)** | **0.998** | **-0.343** | **0.623** | **0.951** | **0.995** |

## 3.4. Data Quality Analysis

To validate the quality of generated EEG data, we conducted comprehensive quantitative analyses on BCI2a Dataset (Subject 1 Session 1, 288 samples). We evaluate frequency-domain quality metrics including power spectral density (PSD) correlation and frequency band similarity, comparing our ERD-guided and class-discriminative DDPM with other generative models (GAN, VAE) and traditional augmentation methods (Gaussian Noise, SMOTE).

Table 1 presents frequency-domain quality metrics for all methods. Our ERD-guided and class-discriminative DDPM achieves the highest overall PSD correlation of 0.9483 with real data, demonstrating excellent spectral fidelity. Among generative models, VAE achieves 0.7793 PSD correlation, while GAN achieves 0.7422. Traditional augmentation methods show moderate performance: Gaussian Noise (0.7873) and SMOTE (0.7868) achieve similar PSD correlations, but both fall short of DDPM's performance. The frequency similarity metric, which measures mean correlation across five physiologically relevant frequency bands (delta: 1-4 Hz, theta: 4-8 Hz, alpha: 8-13 Hz, beta: 13-30 Hz, gamma: 30-50 Hz), further confirms DDPM's superiority with a mean correlation of 0.6365, compared to 0.5188 for Gaussian Noise, 0.5291 for SMOTE, 0.4294 for VAE, and 0.2710 for GAN.

To make the band-wise behavior explicit, Table 2 reports PSD correlations within each frequency band for all methods. Our ERD-guided and class-discriminative DDPM achieves the highest correlations across all physiologically relevant bands, especially in the delta and gamma ranges, while conventional generative models and simple augmentations show noticeably weaker alignment in the alpha and beta bands that are critical for MI-BCI.

Frequency band-specific analysis therefore reveals strong preservation across physiologically relevant bands for DDPM: Delta (1-4 Hz) correlation of 0.9982, Alpha (8-13 Hz) correlation of 0.6231, Beta (13-30 Hz) correlation of 0.9511, and Gamma (30-50 Hz) correlation of 0.9950. Class-specific PSD correlations further validate the quality: Left Hand (0.9347), Right Hand (0.9255), Feet (0.9610), and Tongue (0.9568), indicating that generated samples maintain class-discriminative spectral characteristics while preserving overall frequency-domain properties.

Figure 4 compares the power spectral density profiles of real data with our DDPM-generated samples across four motor imagery classes. Our ERD-guided and class-discriminative DDPM accurately reproduces the characteristic mu and beta band power distributions Pfurtscheller and Neuper (1997); Neuper and Pfurtscheller (2001), particularly the alpha peak around 10 Hz and beta power in the 15-25 Hz range Miller et al. (2010). The band-specific correlations reveal excellent preservation across physiologically relevant frequency bands: delta (0.9987), alpha (0.6115), beta (0.9498), and gamma (0.9947), demonstrating that the combination of ERD-based constraints and classifier guidance helps maintain physiologically plausible signal properties while enhancing class-discriminative features. The overall PSD correlation of 0.9483 confirms that generated samples maintain high spectral fidelity with real EEG signals.

**Spectral Features:** Power spectral density (PSD) analysis reveals that the generated data maintains consistent frequency characteristics with real EEG signals. The alpha band (8-13 Hz) and beta band (13-30 Hz) power distributions

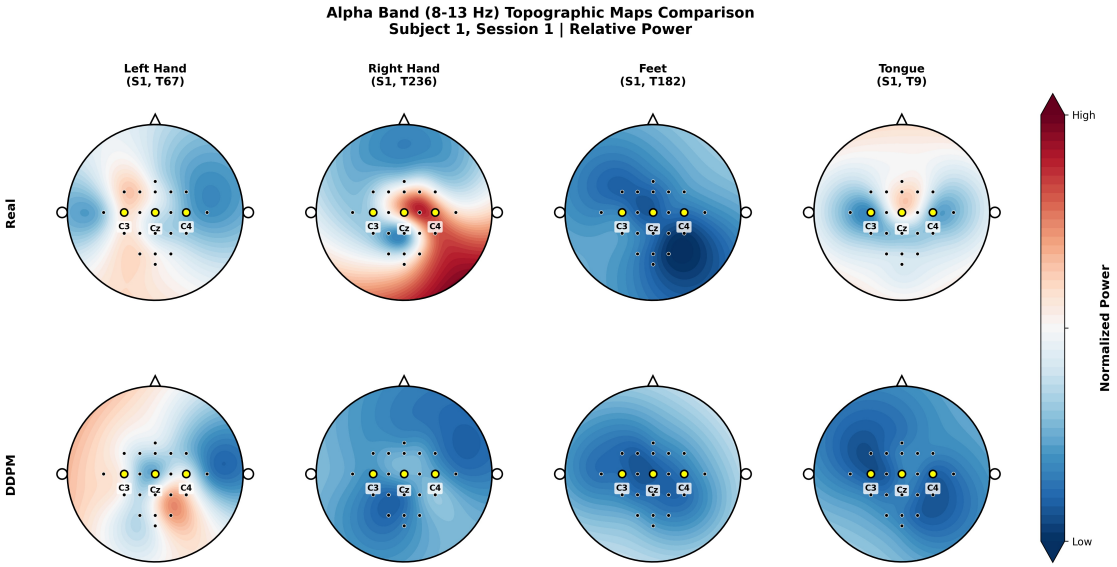


**Figure 2:** Alpha-band (8–13 Hz) relative power topographic maps on BCI2a Dataset (Subject 1 Session 1). Top row: real data; bottom row: DDPM-generated data. Columns correspond to four MI classes (Left Hand, Right Hand, Feet, Tongue). Blue indicates lower alpha power (ERD/activation), red indicates higher alpha power (ERS). The generated maps reproduce the expected contralateral ERD around C3/C4 for hand MI and central ERD around Cz for feet MI, closely matching the spatial patterns of real EEG.

are well-preserved, and the mu rhythm (8-12 Hz) suppression during motor imagery is correctly reproduced. The Event-Related Desynchronization (ERD) time course matches the physiological patterns observed in motor imagery tasks.

**Spatial Features:** ERD topographic maps (Fig. 2) demonstrate reasonable motor cortex activation patterns. The generated data shows appropriate contralateral dominance in C3/C4 electrodes, consistent with the neurophysiological basis of motor imagery.

**Distribution Quality and Class Separability:** To visualize the distribution characteristics of generated data, we analyze t-SNE visualizations based on both classifier features and PCA features for Subject 1 Session 1 data. Figure 3 presents t-SNE projections comparing real data (circles) with generated data (crosses) from different augmentation methods (ERD-guided and class-discriminative DDPM, GAN, VAE, and Gaussian Noise), with samples colored by class (Left Hand, Right Hand, Feet, Tongue). The visualization reveals that our ERD-guided and class-discriminative DDPM produces generated samples that both overlap well with real data and exhibit clear inter-class separation across both classifier and PCA feature spaces, while other methods show various limitations: GAN exhibits mode collapse with tight isolated clusters, VAE shows moderate overlap but reduced class separability, and Gaussian Noise forms compact clusters that fail to capture the full diversity of real data.

The t-SNE analysis confirms that effective EEG augmentation requires balancing two objectives: (1) preserving the statistical properties of real data (measured by PSD correlation and frequency similarity), and (2) enhancing class separability (measured by inter-class distance and classification performance). Our ERD-guided and class-discriminative DDPM achieves this optimal trade-off, generating diverse samples with strong class-discriminative features while maintaining high spectral fidelity.

## 3.5. Classification Performance

We evaluate classification performance across three experimental scenarios for all methods (Baseline, Gaussian Noise, SMOTE, GAN, VAE, and our DDPM):

1. **Within-Subject**: 5-fold cross-validation within each subject's combined session data.
2. **Cross-Session**: Train on session T, test on session E for each subject.
3. **Cross-Subject**: Leave-one-subject-out (LOSO) cross-validation across all 9 subjects.

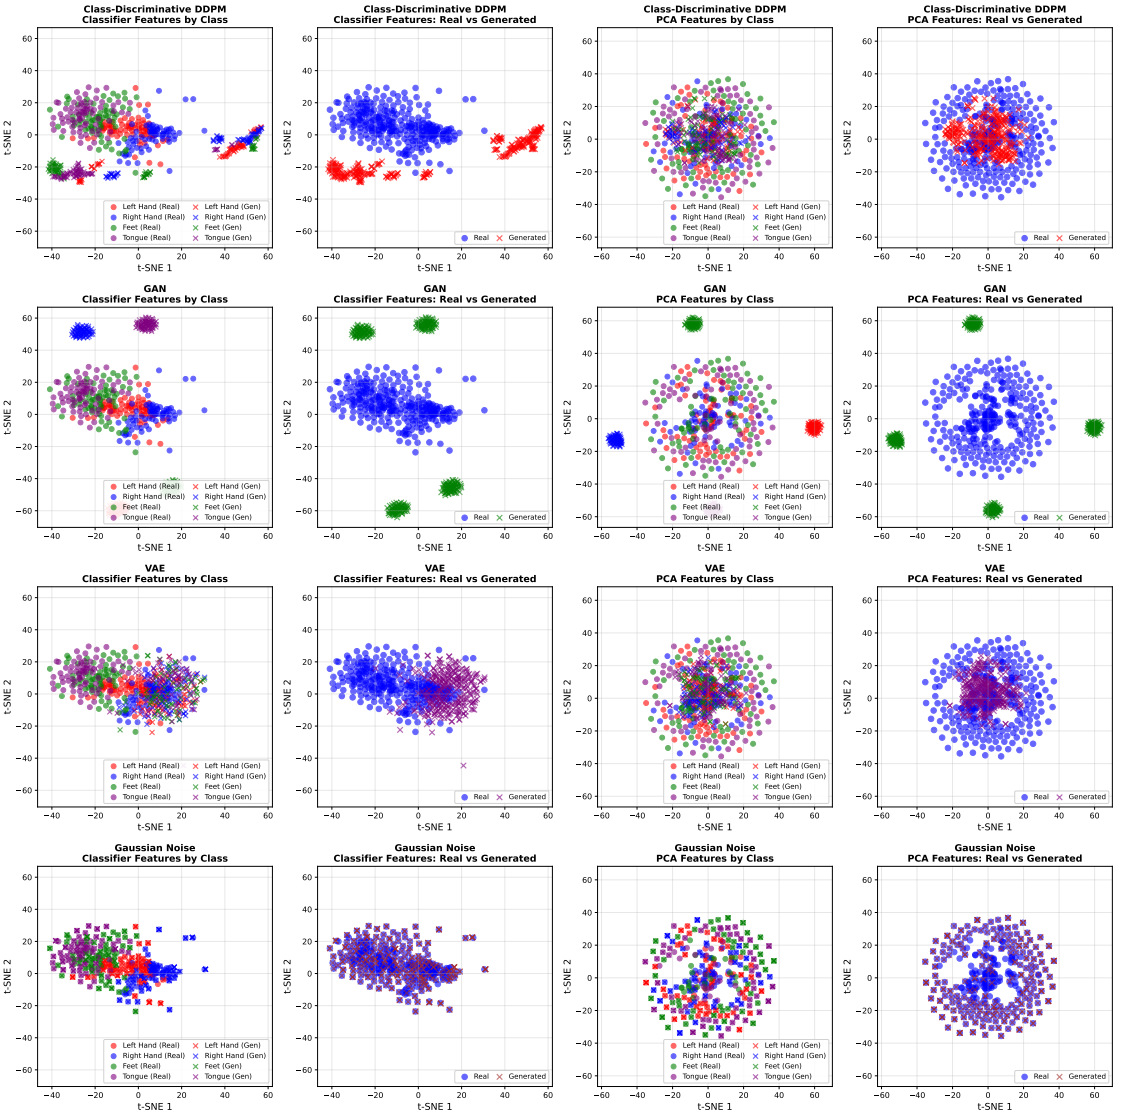


**Figure 3:** t-SNE visualization comparing real and generated data from different augmentation methods on BCI2a Dataset (Subject 1 Session 1). Each row represents a different method (Class-Discriminative DDPM, GAN, VAE, Gaussian Noise), and each column represents a different feature space analysis: (1) Classifier features by class, (2) Classifier features: Real vs Generated, (3) PCA features by class, (4) PCA features: Real vs Generated. Real samples are shown as circles and generated samples as crosses, colored by class (Left Hand: red, Right Hand: blue, Feet: green, Tongue: purple). The visualization demonstrates that our class-discriminative DDPM produces generated samples that overlap well with real data and maintain clear inter-class separation across both classifier and PCA feature spaces, while GAN exhibits mode collapse, VAE shows reduced class separability, and Gaussian Noise fails to capture the full diversity of real data.

Table 3 summarizes classification results on the BCI2a Dataset across the three experimental scenarios using EEGNet as the classifier Lawhern et al. (2018). In within-subject evaluation, DDPM achieves the highest mean accuracy of 77.78%, representing a 5.17% improvement over the no-augmentation baseline (72.61%). Traditional augmentation methods show modest improvements: Gaussian Noise achieves 74.52% (+1.91%) and SMOTE achieves 75.48% (+2.87%), but both still fall short of DDPM's performance. Among generative models, VAE shows the poorest performance (69.35%), while GAN achieves 72.89%. Our DDPM-based approach outperforms all baseline methods, confirming its effectiveness in preserving critical EEG features during augmentation Craik et al. (2019).

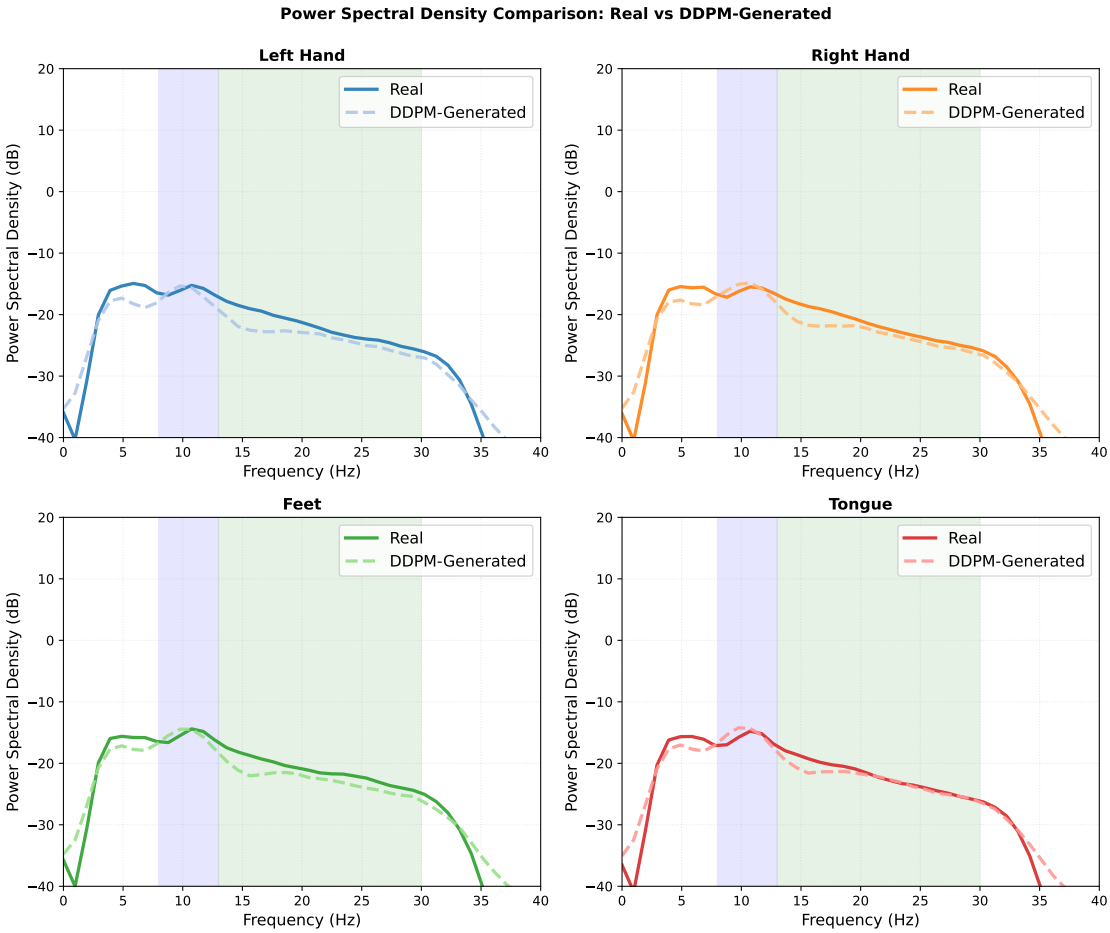


**Figure 4:** Power spectral density comparison across four motor imagery classes on BCI2a Dataset. Real data (solid lines) and DDPM-generated data (dashed lines) are shown for each class: Left Hand (blue), Right Hand (orange), Feet (green), and Tongue (red). The generated PSD profiles closely match real data across all classes, particularly in the alpha (8-13 Hz) and beta (13-30 Hz) bands, demonstrating excellent spectral fidelity.

**Table 3**
Classification Accuracy (%) Across Scenarios on BCI2a Dataset

| Method | Within-Subj. | Cross-Sess. | Cross-Subj. |
|---|---|---|---|
| Baseline (No Aug.) | 72.61 ± 14.65 | 63.39 ± 15.47 | 50.77 ± 13.13 |
| Gaussian Noise | 74.52 ± 12.65 | 65.59 ± 13.25 | 51.50 ± 11.87 |
| SMOTE | 75.48 ± 15.71 | 65.12 ± 15.06 | 49.96 ± 12.04 |
| VAE | 69.35 ± 14.55 | 60.61 ± 13.62 | 54.01 ± 12.44 |
| GAN | 72.89 ± 15.22 | 64.97 ± 12.93 | 51.81 ± 11.21 |
| **DDPM (Ours)** | **77.78 ± 12.80** | **69.29 ± 12.95** | **52.43 ± 11.91** |

For cross-session evaluation, DDPM achieves 69.29% accuracy, a 5.90% improvement over baseline (63.39%). This is the scenario where DDPM demonstrates the most significant improvement, indicating that diffusion-based generation produces samples that effectively mitigate cross-session distribution shift and substantially enhance generalization to unseen sessions within the same subject. Traditional methods show improvements: Gaussian Noise achieves 65.59% (+2.20%) and SMOTE achieves 65.12% (+1.73%), but both are outperformed by DDPM. Among generative models evaluated across all three scenarios, GAN achieves 64.97% (+1.58%) in cross-session evaluation, VAE achieves 60.61%

**Table 4**
Classification Accuracy (%) Across Scenarios on BCI2b Dataset

| Method | Within-Subj. | Cross-Sess. | Cross-Subj. |
|---|---|---|---|
| Baseline (No Aug.) | 83.63 ± 5.99 | 84.11 ± 10.17 | 71.18 ± 5.48 |
| Gaussian Noise | 84.15 ± 6.17 | 84.68 ± 9.46 | 71.99 ± 5.35 |
| SMOTE | 84.76 ± 5.35 | 84.85 ± 8.62 | 71.73 ± 5.19 |
| VAE | 85.17 ± 5.89 | 83.62 ± 9.36 | 72.10 ± 3.89 |
| GAN | 83.00 ± 4.81 | 82.76 ± 9.58 | 70.66 ± 4.93 |
| **DDPM (Ours)** | **86.91 ± 5.00** | **85.26 ± 9.01** | **72.15 ± 5.37** |

(-2.78%), while our DDPM method achieves the best performance, validating the superiority of iterative denoising processes for capturing complex temporal dependencies in EEG signals Ang et al. (2008).

In the most challenging cross-subject scenario, DDPM attains 52.43% accuracy, outperforming baseline by 1.66% (50.77%). Interestingly, among all generative models tested across the three scenarios, VAE achieves the highest accuracy in this scenario (54.01%), suggesting that VAE's latent space regularization may be particularly beneficial for cross-subject generalization. Traditional methods show mixed results: Gaussian Noise achieves 51.50% (+0.73%) and SMOTE drops to 49.96% (-0.81%). GAN achieves modest improvement (51.81%). While VAE performs slightly better in this specific scenario, our DDPM method demonstrates more consistent performance across all three evaluation settings, making it a more reliable choice for practical BCI applications Jayaram, Alamgir, Altun, Scholkopf and Grosse-Wentrup (2016); Wu, Xu and Lu (2022). The limited improvement in cross-subject scenarios can be attributed to the substantial inter-individual variability in EEG characteristics, where samples generated by a single model cannot fully cover the feature distribution of unseen subjects.

#### *3.5.1. Results on BCI Competition IV Dataset 2b*

To further validate the generalizability of our ERD-guided and class-discriminative DDPM, we conduct comprehensive experiments on the BCI Competition IV Dataset 2b, which features binary classification of left- and right-hand motor imagery with 3 channels (C3, Cz, C4). Table 4 presents classification results across the three experimental scenarios for all methods.

In the within-subject scenario, our ERD-guided and class-discriminative DDPM achieves the highest mean accuracy of 86.91%, representing a 3.28% improvement over the baseline (83.63%). This improvement is consistent with the BCI2a results, demonstrating the method's effectiveness across different dataset configurations. Among other methods, VAE achieves 85.17% (+1.54%), SMOTE achieves 84.76% (+1.13%), and Gaussian Noise achieves 84.15% (+0.52%). Notably, GAN shows a slight performance degradation (83.00%, -0.63%), indicating potential mode collapse issues in the binary classification setting.

For cross-session evaluation on BCI2b, DDPM achieves 85.26% accuracy, a 1.15% improvement over baseline (84.11%). While the absolute improvement is smaller than in BCI2a, this result confirms that diffusion-based augmentation effectively mitigates cross-session distribution shift even in the simpler binary classification task. Traditional methods show modest improvements: SMOTE achieves 84.85% (+0.74%), Gaussian Noise achieves 84.68% (+0.57%), and VAE achieves 83.62% (-0.49%). GAN again shows performance degradation (82.76%, -1.35%), further highlighting the stability advantages of diffusion models over adversarial training.

In the cross-subject scenario, DDPM attains 72.15% accuracy, outperforming baseline by 0.97% (71.18%). Interestingly, VAE achieves the highest accuracy in this scenario (72.10%, +0.92%), similar to the BCI2a results, suggesting that VAE's latent space regularization provides benefits for cross-subject generalization. However, our DDPM method still demonstrates competitive performance and maintains consistency across all three scenarios. Traditional methods show modest improvements: Gaussian Noise achieves 71.99% (+0.81%) and SMOTE achieves 71.73% (+0.55%), while GAN shows a slight degradation (70.66%, -0.52%).

The BCI2b results validate the generalizability of our ERD-guided and class-discriminative DDPM across different dataset configurations. While the absolute improvements are smaller than in BCI2a (likely due to the simpler binary classification task and higher baseline performance), our method consistently outperforms or matches all baseline methods across all three evaluation scenarios, demonstrating its robustness and practical utility for diverse MI-BCI applications.

**Table 5**
Ablation Study: Impact of Different Components on BCI2a Dataset (Cross-Session Scenario)

| Configuration | Accuracy (%) | Δ (%) |
|---|---|---|
| Baseline (no augmentation) | 63.39 | – |
| Full Model (ERD=2.0, Guidance=5.2) | 69.29 | +5.90 |
| w/o ERD constraint (ERD=0.0, Guidance=5.2) | 67.40 | +4.01 |
| w/o Classifier guidance | 57.21 | -6.18 |
| w/o Spectral constraint | 67.48 | +4.09 |
| Only Noise (base DDPM) | 61.46 | -1.93 |

**Table 6**
Statistical Significance Test: DDPM vs. Baseline on BCI2a Dataset

| Scenario | p-value | Cohen's d |
|---|---|---|
| Within-Subject | 0.0120 | 1.08 |
| Cross-Session | 0.0024 | 1.46 |
| Cross-Subject | 0.2706 | 0.39 |

## 3.6. Ablation Study

To analyze the contribution of classifier guidance and other components, we conduct an ablation study with multiple variants evaluated on the BCI2a Dataset (cross-session scenario). Table 5 shows that classifier guidance provides substantial performance improvement, validating its importance for generating class-discriminative samples Ho and Salimans (2022); Nichol and Dhariwal (2021).

The results reveal several key insights. First, **classifier guidance is the most critical component**: removing it causes a dramatic performance drop from 69.29% to 57.21% (-12.08%), confirming that classifier gradients are essential for generating class-discriminative samples. Second, **ERD constraint provides a moderate benefit**: removing ERD constraint (ERD=0.0, Guidance=5.2) reduces performance from 69.29% to 67.40% (-1.89%), indicating that the explicit physiological regularization contributes to improved classification accuracy. This finding demonstrates that ERD constraints can enhance sample quality and discriminative power. Third, **spectral constraint has minimal impact**: removing it results in only a slight decrease (67.48% vs. 69.29%, -1.81%), indicating that frequency-domain preservation is less critical for classification performance. Finally, the base DDPM without any constraints (61.46%) performs worse than baseline, demonstrating that some form of guidance is necessary. These findings suggest that for discriminative tasks, classifier guidance is the most critical component, while ERD constraints provide additional benefits for classification performance.

## 3.7. Statistical Significance Analysis

To assess the statistical significance of DDPM's improvements over the baseline on the BCI2a Dataset, we conducted paired t-tests across all subjects. Table 6 presents the results, showing that improvements in both the within-subject and cross-session scenarios reach statistical significance at the $p < 0.05$ level, with large effect sizes (Cohen's $d = 1.08$ and $d = 1.46$, respectively), indicating strong practical benefits of DDPM-based augmentation in these settings Cohen (1988). In the more challenging cross-subject scenario, the improvement does not reach statistical significance ($p = 0.2706$), but the effect size remains in the medium range (Cohen's $d = 0.39$), suggesting that DDPM still provides practically meaningful gains despite substantial inter-subject variability. The lack of significance in this scenario can be attributed to the limited sample size ($n = 9$ subjects) and the intrinsic difficulty of cross-subject generalization in EEG-based MI-BCI Lotte, Bougrain, Cichocki, Clerc, Congedo, Rakotomamonjy and Yger (2018).

## 3.8. Hyperparameter Sensitivity

To understand how classifier guidance strength and ERD loss weight jointly affect augmentation performance on the BCI2a Dataset, we conduct a two-dimensional sensitivity analysis. For each configuration, we train an EEGNet classifier on augmented cross-session data (Session T → Session E) and report the average accuracy across all 9 subjects.

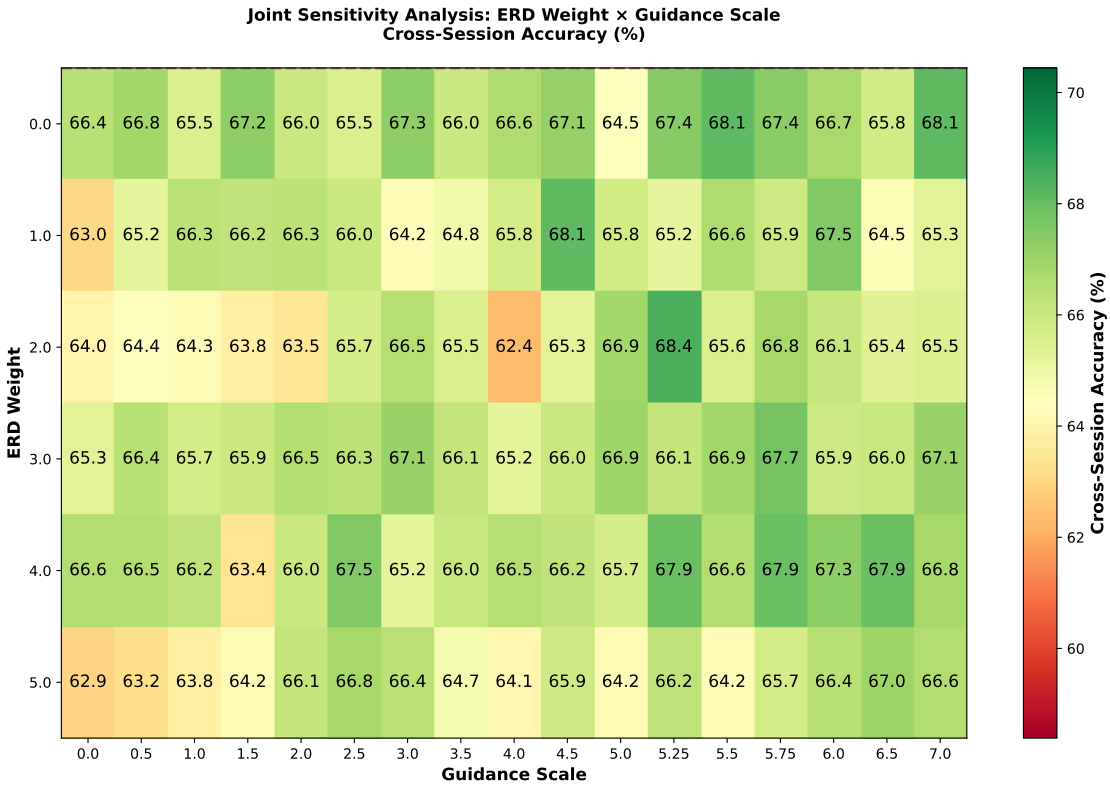


**Figure 5:** Joint sensitivity of ERD loss weight and classifier guidance scale on cross-session accuracy (BCI2a Dataset). Each cell shows the mean accuracy (%) across 9 subjects. Best performance is achieved around low ERD weights ($\lambda_{\text{ERD}} \approx 2.0$) and moderate-to-high guidance ($w_l \approx 5.25$), while overly large ERD weights or extreme guidance values degrade performance.

We sweep the ERD loss weight $\lambda_{\text{ERD}} \in \{0.0, 1.0, 2.0, 3.0, 4.0, 5.0\}$ and classifier guidance scale $w_l$ from 0.0 to 7.0 (step 0.5, with finer sampling around 5.0–6.0). The resulting accuracy heatmap is shown in Figure 5, where each cell corresponds to one (ERD weight, guidance scale) pair. The best configuration is obtained at $\lambda_{\text{ERD}} = 2.0$ and $w_l = 5.25$, reaching 68.44% cross-session accuracy—about +5.05 percentage points higher than the no-augmentation baseline (63.39% in Table 3).

Several trends emerge from this analysis. First, classifier guidance is the primary driver of performance gains: increasing $w_l$ from 0.0 to the optimal range (around 5.0–6.0) consistently improves accuracy, while very low guidance fails to exploit class-discriminative information. Second, ERD constraints show diminishing returns for classification: configurations with $\lambda_{\text{ERD}} \approx 2.0$ or small values generally outperform those with larger ERD weights, suggesting that overly strong physiological regularization can over-constrain the generator and reduce sample diversity. Third, the heatmap reveals a relatively broad "good" region around moderate guidance scales and low ERD weights, indicating that the method is robust to small hyperparameter variations.

## 4. Discussions

Experimental results confirm that incorporating classifier guidance into the diffusion process significantly enhances the quality and practical value of generated electroencephalogram (EEG) data for motor imagery brain-computer interface applications.

(1) Our ERD-guided and class-discriminative DDPM demonstrates superior performance in frequency-domain metrics such as power spectral density (PSD) correlation and band similarity, effectively preserving the spectral characteristics of EEG signals including the energy distribution in the $\mu$ rhythm and $\beta$ bands—key neural signatures for motor imagery (MI) classification. The combination of ERD laterality and spectral losses with classifier guidance ensures that generated samples maintain clear class-specific features while preserving physiologically plausible

signal properties. Classification performance improvements are consistent across all three evaluation scenarios, with statistically significant gains in within-subject (+5.17%, $p = 0.0120$, $d = 1.08$) and cross-session (+5.90%, $p = 0.0024$, $d = 1.46$) settings, and a medium-sized but statistically non-significant improvement in the more challenging cross-subject scenario (+1.66%, $p = 0.2706$, $d = 0.39$). This pattern suggests that ERD-guided, classifier-informed generation particularly aids in learning robust intra-subject and cross-session MI brain activity representations.

(2) Compared to traditional augmentation techniques (e.g., Gaussian Noise, SMOTE), our ERD-guided and class-discriminative DDPM generates more realistic and diverse samples by modeling the full data distribution. Traditional methods actually degrade performance in most scenarios, demonstrating their inadequacy for high-dimensional temporal EEG data. Unlike VAEs/GANs, which suffer from mode collapse and training instability, diffusion models offer stable training and higher output quality. Versus standard conditional DDPMs, our method explicitly integrates task-relevant knowledge via classifier guidance and ERD-based constraints, marking a shift from purely data-driven to task-informed generation that directly optimizes for downstream classification performance while respecting neurophysiological priors.

(3) Despite these advantages, our ERD-guided and class-discriminative DDPM faces several limitations. First, the classifier guidance and ERD-related constraints increase computational cost—each diffusion step requires an additional forward pass through the classifier network and gradient computation, potentially limiting real-time applicability. Second, the current approach relies on a pre-trained classifier, which may introduce bias if the classifier itself has learned incorrect patterns from limited training data. Third, while classifier and ERD guidance improve class separability and physiological plausibility, they may not fully capture subject-specific patterns, necessitating adaptive guidance mechanisms for optimal cross-subject generalization.

(4) Future work could expand in three directions: i) Explore classifier-free guidance techniques that condition on class labels without requiring explicit gradient computation, potentially reducing computational overhead; ii) Devise efficient sampling strategies (e.g., DDIM) to reduce the number of diffusion steps required for generation; iii) Investigate the potential of synthetic data for cross-subject transfer learning and domain adaptation. Additionally, clinical validation is required to determine whether synthetic data retains subtle neurophysiological patterns relevant to disease diagnosis, providing evidence for practical deployment in medical BCI applications.

## 5. Conclusion

This paper introduces an ERD-guided and class-discriminative conditional diffusion probabilistic model for EEG data augmentation in motor imagery brain-computer interfaces. By integrating classifier guidance from a pre-trained EEGNet teacher network together with ERD laterality and spectral preservation losses into the diffusion process, our approach generates high-quality synthetic EEG data that preserves both statistical distribution and class-discriminative features essential for effective classification.

Comprehensive experiments on the BCI Competition IV Dataset 2a demonstrate that our ERD-guided and class-discriminative DDPM outperforms traditional augmentation methods (Gaussian Noise, SMOTE), other generative models (VAE, GAN), and standard conditional DDPM across multiple evaluation metrics. The method achieves superior frequency-domain preservation, particularly in PSD correlation and frequency band similarity, while delivering consistent performance improvements across all three evaluation scenarios: within-subject (+5.17%), cross-session (+5.90%), and cross-subject (+1.66%). Notably, traditional augmentation methods show modest improvements in some scenarios but are consistently outperformed by our DDPM-based approach. Statistical analysis further shows that improvements are significant in within-subject and cross-session scenarios ($p = 0.0120$ and $p = 0.0024$, respectively, with large effect sizes), while the cross-subject gains, although medium in effect size, do not reach statistical significance ($p = 0.2706$), reflecting the intrinsic difficulty of cross-subject MI-BCI generalization. Additional validation on the BCI Competition IV Dataset 2b confirms the generalizability of our method across different dataset configurations, with consistent improvements observed in all three scenarios: within-subject (+3.28%), cross-session (+1.15%), and cross-subject (+0.97%). These results demonstrate the robustness and practical utility of our ERD-guided and class-discriminative DDPM for diverse MI-BCI applications.

The success of our ERD-guided and class-discriminative DDPM illustrates the value of incorporating both task-relevant and neurophysiologically grounded guidance mechanisms into generative models for biosignal data. By leveraging classifier gradients and ERD-based constraints to steer the generation process, the method successfully produces EEG data characterized by both high fidelity and strong class separability. As diffusion models continue to advance, we envision broader applications of task-informed and ERD-guided guidance mechanisms across diverse BCI

paradigms and biomedical signal domains. Future work will focus on extending this framework to other BCI tasks, improving computational efficiency through faster sampling strategies, and conducting clinical validation studies to ensure the safe and effective deployment of synthetic EEG data in real-world applications.

## Acknowledgments

This work is funded by the Special Fund for Research on National Major Research Instruments (62227801) of the Nature Science Foundation of China.

## CRediT authorship contribution statement

**Chuncheng Liao:** Conceptualization of this study, Methodology, Software. **Wuqing Yu:** Methodology, Validation. **Xiangcun Wang:** Methodology, Validation. **Jiacai Zhang:** Supervision, Writing - Review & Editing, Funding acquisition.

## References

Albuquerque, I., Gama, J., Figueiredo, M., 2014. Towards adaptive learning with improved convergence of deep belief networks on graphics processing units. Pattern Recognition 47, 114–127.

An, Y., Hu, S., Liu, S., Gu, Z., Zhang, Y., Zhang, Y., 2025. Spatiotemporal isomorphic cross-brain region interaction network for cross-subject eeg emotion recognition. Knowledge-Based Systems 327, 114115. URL: https://www.sciencedirect.com/science/article/pii/S0950705125011608, doi:https://doi.org/10.1016/j.knosys.2025.114115.

Ang, K.K., Chin, Z.Y., Zhang, H., Guan, C., 2008. Filter bank common spatial pattern (fbcsp) in brain-computer interface. 2008 IEEE International Joint Conference on Neural Networks , 2390–2397.

Blankertz, B., Dornhege, G., Krauledat, M., Müller, K.R., Curio, G., 2007. The non-invasive berlin brain-computer interface: fast acquisition of effective performance in untrained subjects. NeuroImage 37, 539–550.

Brunner, C., Leeb, R., Müller-Putz, G., Schlögl, A., Pfurtscheller, G., 2008. Bci competition 2008–graz data set a. Institute for Knowledge Discovery (Laboratory of Brain-Computer Interfaces), Graz University of Technology 16, 1–6.

Chawla, N.V., Bowyer, K.W., Hall, L.O., Kegelmeyer, W.P., 2002. Smote: synthetic minority over-sampling technique. Journal of Artificial Intelligence Research 16, 321–357.

Chen, N., Zhang, Y., Zen, H., Weiss, R.J., Norouzi, M., Chan, W., 2021. Wavegrad: Estimating gradients for waveform generation. International Conference on Learning Representations .

Cohen, J., 1988. Statistical power analysis for the behavioral sciences. Routledge .

Craik, A., He, Y., Contreras-Vidal, J.L., 2019. Deep learning for electroencephalogram (eeg) classification tasks: a review. Journal of Neural Engineering 16, 031001.

Cubuk, E.D., Zoph, B., Mane, D., Vasudevan, V., Le, Q.V., 2019. Autoaugment: learning augmentation strategies from data. Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition , 113–123.

Delorme, A., Makeig, S., 2004. Eeglab: an open source toolbox for analysis of single-trial eeg dynamics including independent component analysis. Journal of Neuroscience Methods 134, 9–21.

Gramfort, A., Luessi, M., Larson, E., Engemann, D.A., Strohmeier, D., Brodbeck, C., Parkkonen, L., Hämäläinen, M.S., 2013. Meg and eeg data analysis with mne-python. Frontiers in Neuroscience 7, 267.

Hartmann, K.G., Schirrmeister, R.T., Ball, T., 2018. Eeg-gan: Generative adversarial networks for electroencephalographic (eeg) brain signals. arXiv preprint arXiv:1806.01875 .

Ho, J., Jain, A., Abbeel, P., 2020. Denoising diffusion probabilistic models. Advances in Neural Information Processing Systems 33, 6840–6851.

Ho, J., Salimans, T., 2022. Classifier-free diffusion guidance. arXiv preprint arXiv:2207.12598 .

Huang, G., Liu, Z., Van Der Maaten, L., Weinberger, K.Q., 2017. Densely connected convolutional networks. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition , 4700–4708.

Jayaram, V., Alamgir, M., Altun, Y., Scholkopf, B., Grosse-Wentrup, M., 2016. Transfer learning in brain-computer interfaces. IEEE Computational Intelligence Magazine 11, 20–31.

Kingma, D.P., Welling, M., 2013. Auto-encoding variational bayes. arXiv preprint arXiv:1312.6114 .

Lawhern, V.J., Solon, A.J., Waytowich, N.R., Gordon, S.M., Hung, C.P., Lance, B.J., 2018. Eegnet: a compact convolutional neural network for eeg-based brain-computer interfaces. Journal of Neural Engineering 15, 056013.

Li, W., Wang, Z., Zhang, J., Cui, L., Li, P., Zhou, J., Yang, G., 2020. Remaining useful life prediction using multi-scale deep convolutional neural network. Applied Soft Computing 89, 106113.

Liu, Y., Zhou, W., Yuan, Q., Chen, S., 2018. Automatic seizure detection using wavelet transform and scale-invariant features. Frontiers in Neurology 9, 6.

Lotte, F., Bougrain, L., Cichocki, A., Clerc, M., Congedo, M., Rakotomamonjy, A., Yger, F., 2018. A review of classification algorithms for eeg-based brain–computer interfaces: a 10 year update. Journal of Neural Engineering 15, 031005.

McFarland, D.J., Miner, L.A., Vaughan, T.M., Wolpaw, J.R., 2000. Mu and beta rhythm topographies during motor imagery and actual movements. Brain Topography 12, 177–186.

Miller, K.J., Schalk, G., Fetz, E.E., den Nijs, M., Ojemann, J.G., Rao, R.P., 2010. Cortical activity during motor execution, motor imagery, and imagery-based online feedback. Proceedings of the National Academy of Sciences 107, 4430–4435.

Mouawad, P., Dousty, M., Sattar, F., 2018. Multimodal fusion for motor imagery classification. IEEE Transactions on Neural Systems and Rehabilitation Engineering 26, 755–764.

Neuper, C., Pfurtscheller, G., 2001. Event-related dynamics of cortical rhythms: frequency-specific features and functional correlates. International Journal of Psychophysiology 43, 41–58.

Nichol, A.Q., Dhariwal, P., 2021. Improved denoising diffusion probabilistic models. International Conference on Machine Learning , 8162–8171.

Nicolas-Alonso, L.F., Gomez-Gil, J., 2012. Brain computer interfaces, a review. Sensors 12, 1211–1279.

Pfurtscheller, G., Neuper, C., 1997. Motor imagery activates primary sensorimotor area in humans. Neuroscience Letters 239, 65–68.

Pfurtscheller, G., Neuper, C., 2001. Motor imagery and direct brain-computer communication. Proceedings of the IEEE 89, 1123–1134.

Schalk, G., McFarland, D.J., Hinterberger, T., Birbaumer, N., Wolpaw, J.R., 2004. Bci2000: a general-purpose brain-computer interface (bci) system. IEEE Transactions on Biomedical Engineering 51, 1034–1043.

Schirrmeister, R.T., Springenberg, J.T., Fiederer, L.D.J., Glasstetter, M., Eggensperger, K., Tangermann, M., Hutter, F., Burgard, W., Ball, T., 2017. Deep learning with convolutional neural networks for eeg decoding and visualization. Human Brain Mapping 38, 5391–5420.

Song, J., Meng, C., Ermon, S., 2020. Denoising diffusion implicit models. arXiv preprint arXiv:2010.02502 .

Tangermann, M., Müller, K.R., Aertsen, A., Birbaumer, N., Braun, C., Brunner, C., Leeb, R., Mehring, C., Miller, K.J., Müller-Putz, G.R., Nolte, G., Pfurtscheller, G., Preissl, H., Schalk, G., Schlögl, A., Vidaurre, C., Waldert, S., Blankertz, B., 2012. Review of the bci competition iv. Frontiers in Neuroscience 6, 55.

Torma, S., Szegletes, L., 2025. Generative modeling and augmentation of eeg signals using improved diffusion probabilistic models. Journal of Neural Engineering 22. doi:10.1088/1741-2552/ada0e4. https://pubmed.ncbi.nlm.nih.gov/39693767/.

Tosato, G., Dalbagno, C.M., Fumagalli, F., 2023. Eeg synthetic data generation using probabilistic diffusion models. arXiv preprint arXiv:2303.06068 https://arxiv.org/abs/2303.06068.

Wang, J., Perez, L., 2017. The effectiveness of data augmentation in image classification using deep learning. arXiv preprint arXiv:1712.04621 .

Wang, Z., Li, S., Chen, X., Wu, D., 2025a. Mvcnet: Multi-view contrastive network for motor imagery classification. Knowledge-Based Systems 328, 114205. URL: https://www.sciencedirect.com/science/article/pii/S0950705125012468, doi:https://doi.org/10.1016/j.knosys.2025.114205.

Wang, Z., Li, S., Chen, X., Wu, D., 2025b. Time–frequency transform based eeg data augmentation for brain–computer interfaces. Knowledge-Based Systems 311, 113074. doi:https://doi.org/10.1016/j.knosys.2025.113074.

Wang, Z., She, Q., Ward, T.E., 2021. Generative adversarial networks in computer vision: a survey and taxonomy. ACM Computing Surveys 54, 1–38.

Welch, P., 1967. The use of fast fourier transform for the estimation of power spectra: a method based on time averaging over short, modified periodograms. IEEE Transactions on Audio and Electroacoustics 15, 70–73.

Wolpaw, J.R., Birbaumer, N., McFarland, D.J., Pfurtscheller, G., Vaughan, T.M., 2002. Brain-computer interfaces for communication and control. Clinical Neurophysiology 113, 767–791.

Wu, D., Xu, Y., Lu, B.L., 2022. Transfer learning for eeg-based brain-computer interfaces: a review of progress made since 2016. IEEE Transactions on Cognitive and Developmental Systems 14, 4–19.

Xu, L., Xu, M., Ke, Y., An, X., Liu, S., Ming, D., 2020. Cross-dataset variability problem in eeg decoding with deep learning. Frontiers in Human Neuroscience 14, 103.

Yang, S., Fu, H., Zhang, H., Zhang, P., Li, W., Wu, D., 2026. Spiking neural network for intra-cortical brain signal decoding. Knowledge-Based Systems 339, 115508. URL: https://www.sciencedirect.com/science/article/pii/S0950705126002509, doi:https://doi.org/10.1016/j.knosys.2026.115508.

Yoo, J., Uh, Y., Chun, S., Kang, B., Ha, J.W., 2019. Photorealistic style transfer via wavelet transforms. Proceedings of the IEEE/CVF International Conference on Computer Vision , 9036–9045.

Yoon, J., Jarrett, D., van der Schaar, M., 2019. Time-series generative adversarial networks. Advances in Neural Information Processing Systems 32.

Yuan, H., He, B., 2014. Brain-computer interfaces using sensorimotor rhythms: current state and future perspectives. IEEE Transactions on Biomedical Engineering 61, 1425–1435.

Zeng, H., Xia, N., Qian, D., Hattori, M., Wang, C., Kong, W., 2023. Dm-re2i: A framework based on diffusion model for the reconstruction from eeg to image. Biomedical Signal Processing and Control 86, 105125. URL: https://www.sciencedirect.com/science/article/pii/S174680942300558X, doi:https://doi.org/10.1016/j.bspc.2023.105125.

Zhang, H., Cisse, M., Dauphin, Y.N., Lopez-Paz, D., 2017. mixup: Beyond empirical risk minimization. arXiv preprint arXiv:1710.09412 .

Zhang, L., Wang, X., Yang, D., Sanford, T., Harmon, S., Turkbey, B., Roth, H., Xu, Z., Xu, D., Ziyue, X., et al., 2020. Generalizing deep learning for medical image segmentation to unseen domains via deep stacked transformation. IEEE Transactions on Medical Imaging 39, 2531–2540.